

```
In [42]: 1 import scipy as sp
2 import scipy.stats as stats
3 import pandas as pd
4 import numpy as np
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 import copy
8 # Set color map to have light blue background
9 sns.set()
10 import statsmodels.formula.api as smf
11 import statsmodels.api as sm
12 from sklearn.model_selection import train_test_split
13 %matplotlib inline
```

Purpose of the Project

The goal of this project is to analyze, through supervised learning, any kinds of relations between monthly average urban Consumer Price Index, price of grocery items, and time in the United States. I will be using singular and multiple linear regression through Ordinary Least Squares to predict causal relationships between individual grocery items, the monthly CPI, and the months and years that each of these attributes are being evaluated.

I will be using data from the Federal Reserve Economic Data (FRED) Site from the St. Louis Federal Reserve Branch, which has more than 100 years of financial data on the average monthly price of bread, milk, etc. The goal of the project is to evaluate what effects that raising the prices of grocery items have become in response to raising the urban Consumer Price Index (CPI) on urban areas in the United States. The data collected and cleaned is based on a rolling monthly average.

About the Data

The datasets are compilations of grocery items' monthly average prices in urban areas of the United States.

There were some empty data points with each of the CSV files that I was working with, but they have been edited to make sure that any empty values were filled or removed accordingly. There were six CSV files combined into one CSV file containing all of the columns of average grocery items' prices. The dates were also reformatted in the CSV files using Microsoft Excel where the dates, initially listed in its respective column in the form **M/D/YYYY** are now listed in the form **YYYY-MM-DD**.

There are seven columns, one of which is a date consisting of a month and year. By default, the day of the month had already been preset by FRED to the first day of each month for 300 months, or 25 years, from February 1999 to February 2024, which was the last month of recorded data on monthly average urban prices.

The grocery items being compared and run through for modeling are: white bread, ground beef, eggs, whole milk, and bananas. Another column, `urban_cpi`, is the average urban Consumer Price Index (CPI) for spending on groceries in urban US areas.

Data Preparation

All data editing and preparation was done prior to evaluating and/or modeling on any of its contents using Microsoft Excel. For each of the grocery items being evaluated and inputted for modeling, collections of monthly average prices began as early as 1985 for items. The monthly average urban CPI had data collected since January 1913.

Since the analysis and modeling are for patterns and trends in the last 25 years, original files were modified and redacted to only present data for all six attributes listed above from February 1999 to February 2024.

Supervised Learning Models Used

In this analysis, I will use linear regression modeling with Ordinary Least-Squares (OLS) Regression. This will be done by using the `statsmodels` library in Python with `ols`.

```
In [43]: 1 # Call the CSV file and read it. Then return its basic structural info.
          2 groceries_df = pd.read_csv("fred_groceries_021999_022024.csv", encoding='utf-8')
          3 groceries_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 7 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   DATE            301 non-null   object
 1   white_bread     301 non-null   float64
 2   ground_beef    301 non-null   float64
 3   egg             301 non-null   float64
 4   whole_milk     301 non-null   float64
 5   banana         301 non-null   float64
 6   urban_cpi       301 non-null   float64
dtypes: float64(6), object(1)
memory usage: 16.6+ KB
```

```
In [44]: 1 # groceries_df.urban_cpi
```

```
In [45]: 1 # Add a column for signifying the number of each month (300 in total)
          2 month = np.arange(0, len(groceries_df.DATE), 1, dtype=int)
          3 # month
          4 groceries_df['month_number'] = month
          5 groceries_df['month_number'].head()
```

```
Out[45]: 0    0
          1    1
          2    2
          3    3
          4    4
          Name: month_number, dtype: int32
```

In [46]:

```
1 # To refresh, these are the descriptive statistics of the table
2 groceries_df.describe()
```

Out[46]:

	white_bread	ground_beef	egg	whole_milk	banana	urban_cpi	month_number
count	301.000000	301.000000	301.000000	301.000000	301.000000	301.000000	301.000000
mean	1.297900	2.999774	1.651575	3.289449	0.561970	227.267920	150.000000
std	0.258624	1.078457	0.579640	0.392196	0.049593	42.223027	87.035433
min	0.878000	0.000000	0.838000	2.656000	0.469000	163.300000	0.000000
25%	1.046000	2.196000	1.257000	2.964000	0.509000	190.400000	75.000000
50%	1.358000	2.818000	1.599000	3.241000	0.574000	229.554000	150.000000
75%	1.419000	3.841000	1.920000	3.557000	0.601000	251.238000	225.000000
max	2.033000	5.353000	4.823000	4.218000	0.643000	327.731000	300.000000

In [47]:

```
1 # Perform single linear regresssion on each of the columns
2 # with respect to Average Urban Consumer Price Index in US
3 # Create X, y variables
4
5 X = sm.add_constant(groceries_df)
6 y = groceries_df['urban_cpi']
7
8 # Create train and test data splits, 80-20 split
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random.
```

In [48]:

```
1 # Get shapes of X, y
2 print(X.shape, y.shape)
```

(301, 9) (301,)

In [49]:

```
1 # Get shapes of X_train, X_test, y_train, y_test
2
3 print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

(240, 9) (61, 9) (240,) (61,)

```
In [50]: 1 # Get the linear regression model summary
2 # using Ordinary Linear Squares (OLS) from statsmodels
3
4 lr_model = smf.ols(formula="urban_cpi ~ white_bread", data=X_train).fit()
5 lr_model.summary()
```

Out[50]: OLS Regression Results

Dep. Variable:	urban_cpi	R-squared:	0.878
Model:	OLS	Adj. R-squared:	0.877
Method:	Least Squares	F-statistic:	1705.
Date:	Tue, 30 Apr 2024	Prob (F-statistic):	1.69e-110
Time:	10:36:45	Log-Likelihood:	-991.25
No. Observations:	240	AIC:	1986.
Df Residuals:	238	BIC:	1993.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	29.1750	4.895	5.961	0.000	19.533	38.817
white_bread	152.2903	3.688	41.290	0.000	145.024	159.556

Omnibus:	8.429	Durbin-Watson:	2.159
Prob(Omnibus):	0.015	Jarque-Bera (JB):	8.079
Skew:	0.400	Prob(JB):	0.0176
Kurtosis:	2.590	Cond. No.	10.3

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

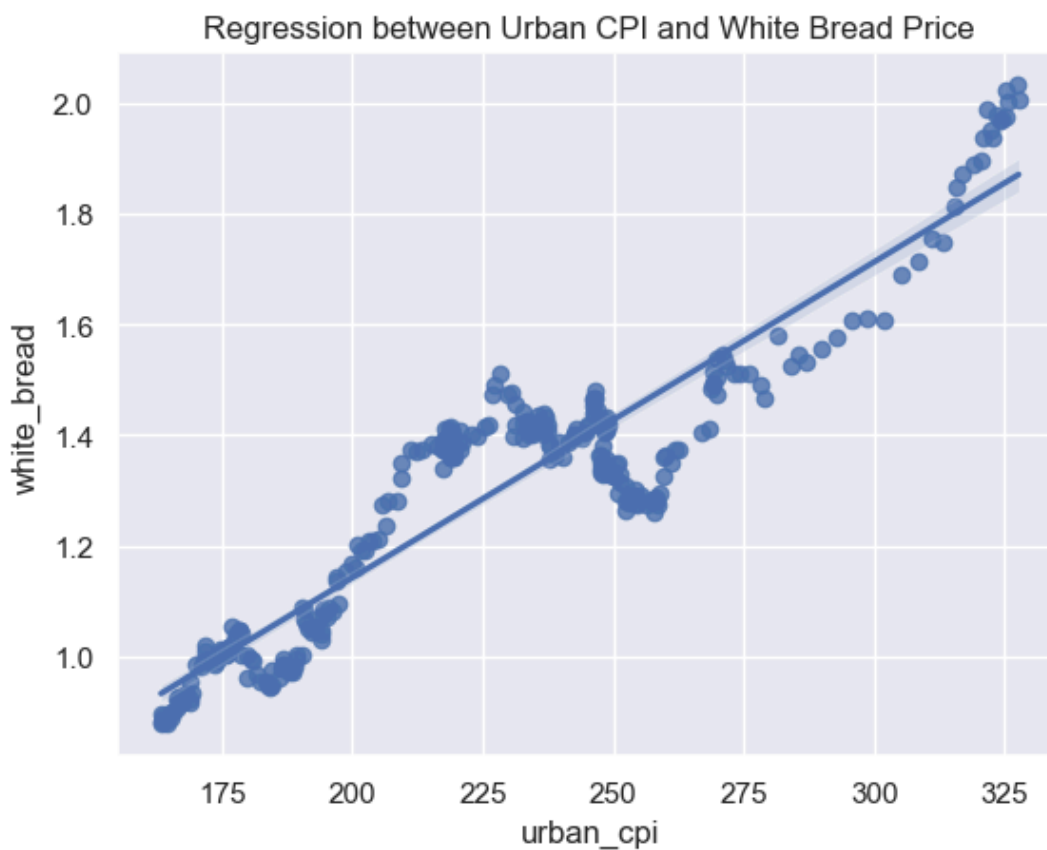
```
In [51]: 1 # Predict the values of CPI on white bread price
2 # then compare to the actual CPI values
3
4 lr_preds = lr_model.predict(X_test)
5 lr_preds
```

```
Out[51]: 223    234.614584
150    253.650872
226    229.589004
296    334.060152
52     175.525946
...
137    236.289777
227    224.258843
26     182.683590
106    224.106553
92     196.237427
Length: 61, dtype: float64
```

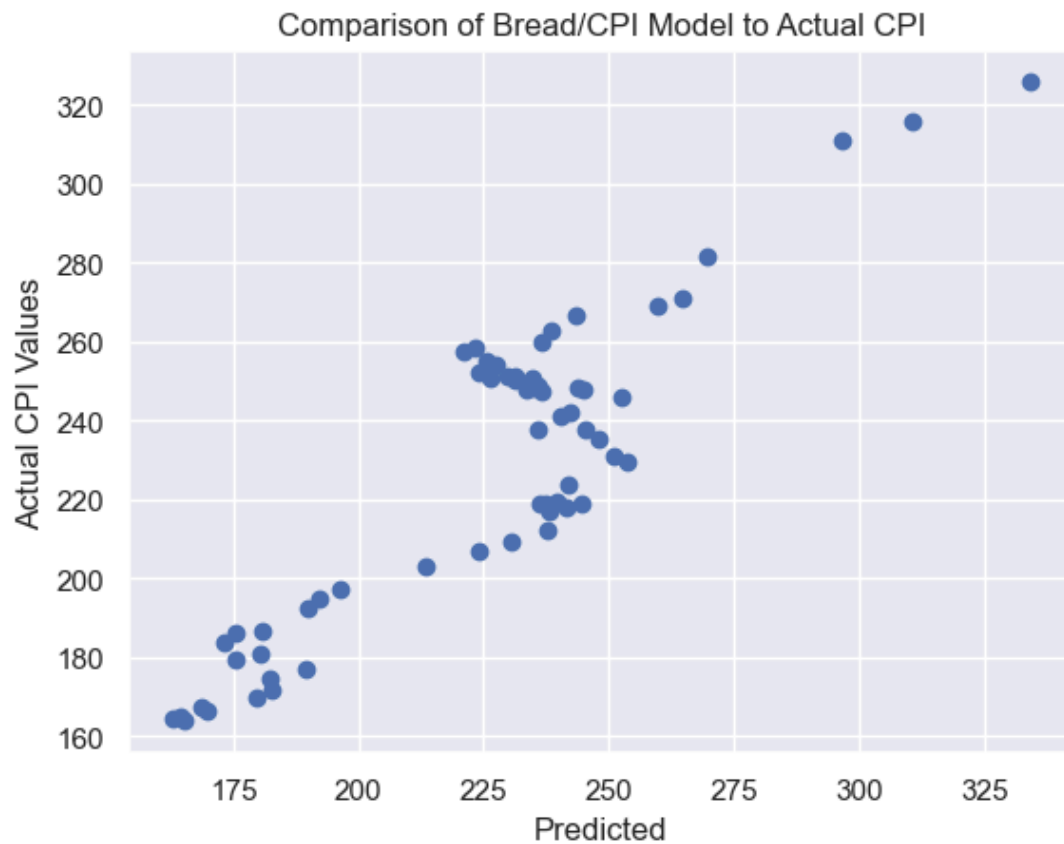
```
In [52]: 1 lr_model.fittedvalues
```

```
Out[52]: 74      182.074429
239      223.192811
64       178.267172
294      329.186863
286      314.414703
...
251      234.919165
192      247.711550
117      238.878713
47       187.861461
172      248.320711
Length: 240, dtype: float64
```

```
In [53]: 1 # Regression plot between urban CPI and white bread price
2
3 sns.regplot(x='urban_cpi', y='white_bread', data=groceries_df)
4 plt.title("Regression between Urban CPI and White Bread Price")
5 plt.show()
```



```
In [56]: 1 # Make a scatter plot of the predicted CPI values (x) compared to the actual list
2
3 plt.scatter(lr_preds, y_test)
4 plt.title("Comparison of Bread/CPI Model to Actual CPI")
5 plt.xlabel("Predicted")
6 plt.ylabel("Actual CPI Values")
7 plt.show()
```



```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [14]: 1 # Single linear regression (OLS) between urban CPI and price of ground beef on X_train
          2
          3 lr_model_beef = smf.ols(formula="urban_cpi ~ ground_beef", data=X_train).fit()
          4 lr_model_beef.summary()
```

Out[14]: OLS Regression Results

Dep. Variable:	urban_cpi	R-squared:	0.895
Model:	OLS	Adj. R-squared:	0.895
Method:	Least Squares	F-statistic:	2038.
Date:	Tue, 30 Apr 2024	Prob (F-statistic):	1.11e-118
Time:	10:34:49	Log-Likelihood:	-972.26
No. Observations:	240	AIC:	1949.
Df Residuals:	238	BIC:	1955.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	116.0974	2.621	44.289	0.000	110.933	121.261
ground_beef	37.0788	0.821	45.144	0.000	35.461	38.697

Omnibus:	182.554	Durbin-Watson:	1.982
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4529.403
Skew:	2.633	Prob(JB):	0.00
Kurtosis:	23.621	Cond. No.	10.1

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

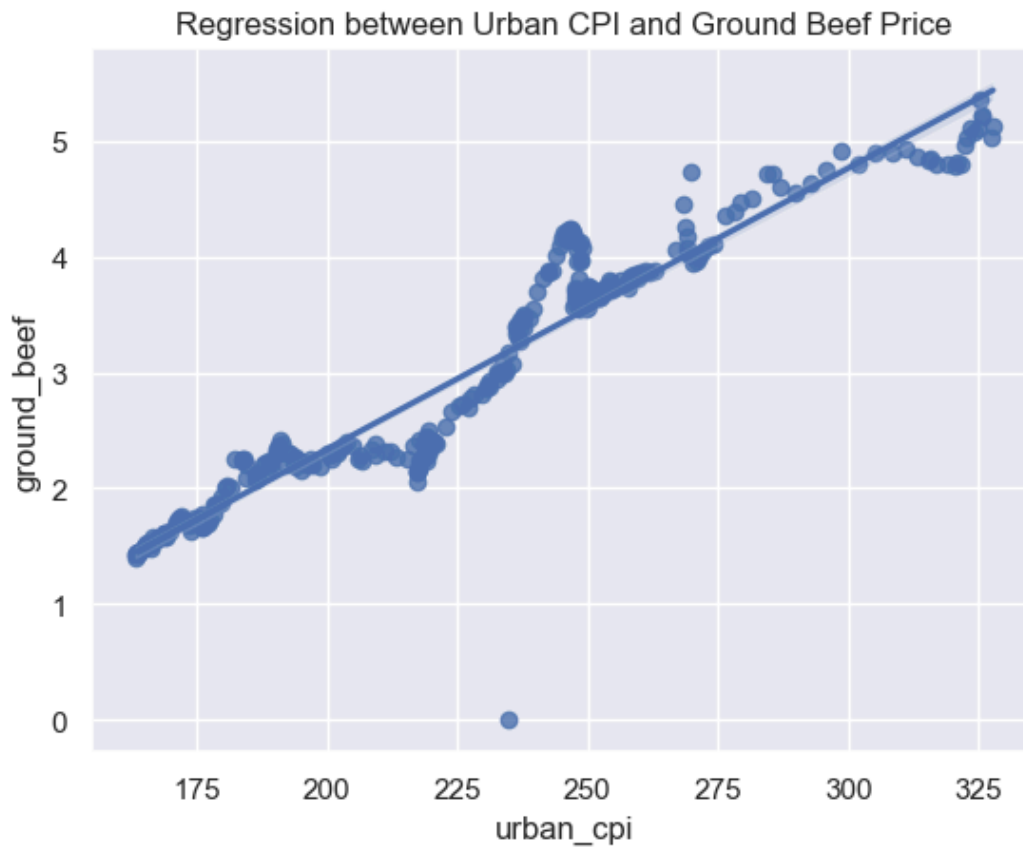
```
In [15]: 1 lr_preds_beef = lr_model_beef.predict(X_test)
          2 lr_preds_beef
```

Out[15]: 223 253.548465
150 220.622491
226 253.585544
296 309.871161
52 187.733596
...
137 207.051651
227 251.101264
26 181.467279
106 198.931394
92 197.893187
Length: 61, dtype: float64

```
In [16]: 1 lr_model_beef.fittedvalues
```

```
Out[16]: 74      201.712304
239      257.033872
64       193.147101
294      304.309341
286      294.075593
...
251      260.185570
192      273.237308
117      203.492086
47       179.947048
172      241.497855
Length: 240, dtype: float64
```

```
In [17]: 1 sns.regplot(x='urban_cpi', y='ground_beef', data=groceries_df)
2 plt.title("Regression between Urban CPI and Ground Beef Price")
3 plt.show()
```



```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```



```
In [18]: 1 lr_model_egg = smf.ols(formula="urban_cpi ~ egg", data=X_train).fit()
          2 lr_model_egg.summary()
```

Out[18]: OLS Regression Results

Dep. Variable:	urban_cpi	R-squared:	0.551
Model:	OLS	Adj. R-squared:	0.549
Method:	Least Squares	F-statistic:	291.8
Date:	Tue, 30 Apr 2024	Prob (F-statistic):	3.08e-43
Time:	10:34:49	Log-Likelihood:	-1147.2
No. Observations:	240	AIC:	2298.
Df Residuals:	238	BIC:	2305.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	138.4136	5.525	25.053	0.000	127.530	149.297
egg	53.4937	3.132	17.081	0.000	47.324	59.663

Omnibus:	19.382	Durbin-Watson:	2.101
Prob(Omnibus):	0.000	Jarque-Bera (JB):	22.551
Skew:	0.748	Prob(JB):	1.27e-05
Kurtosis:	2.869	Cond. No.	6.75

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [19]: 1 lr_preds_egg = lr_model_egg.predict(X_test)
          2 lr_preds_egg
```

Out[19]:

```
223    214.481638
150    229.941320
226    235.504666
296    249.252550
52     202.552540
...
137    215.498018
227    233.043955
26     193.405116
106    250.696880
92     205.655175
Length: 61, dtype: float64
```

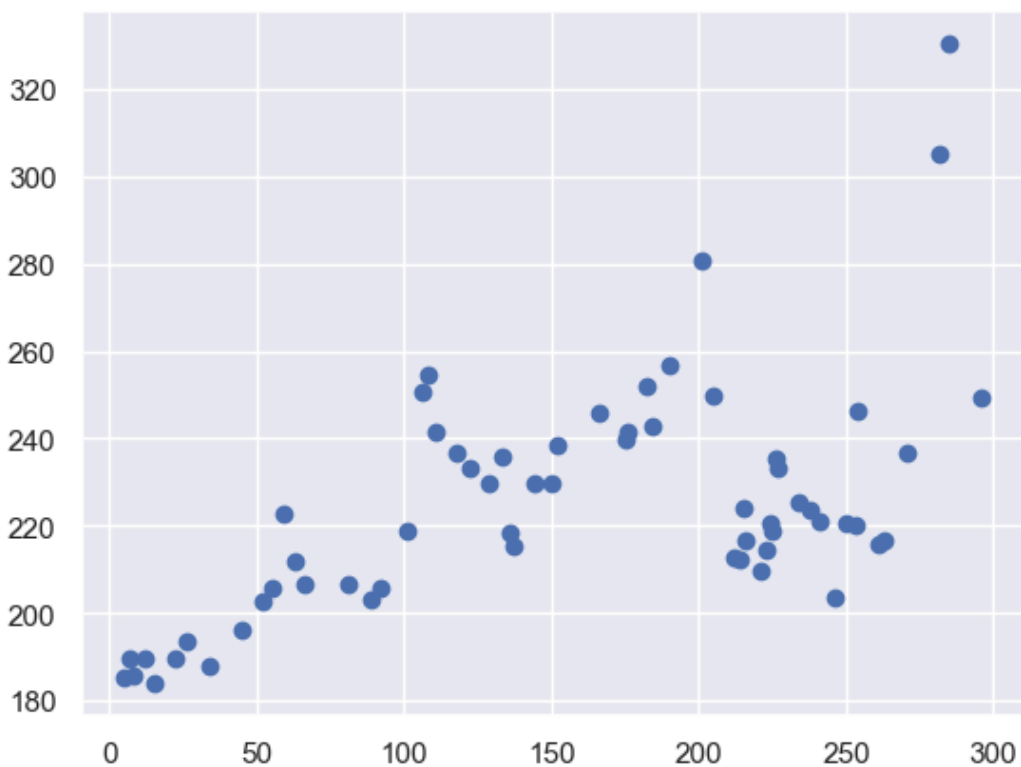
```
In [20]: 1 lr_model_egg.fittedvalues
```

```
Out[20]: 74      200.680260
          239      221.542808
          64      208.543836
          294      247.701232
          286      365.761852
          ...
          251      216.567892
          192      250.108449
          117      236.735021
          47       201.268691
          172      237.858389
          Length: 240, dtype: float64
```

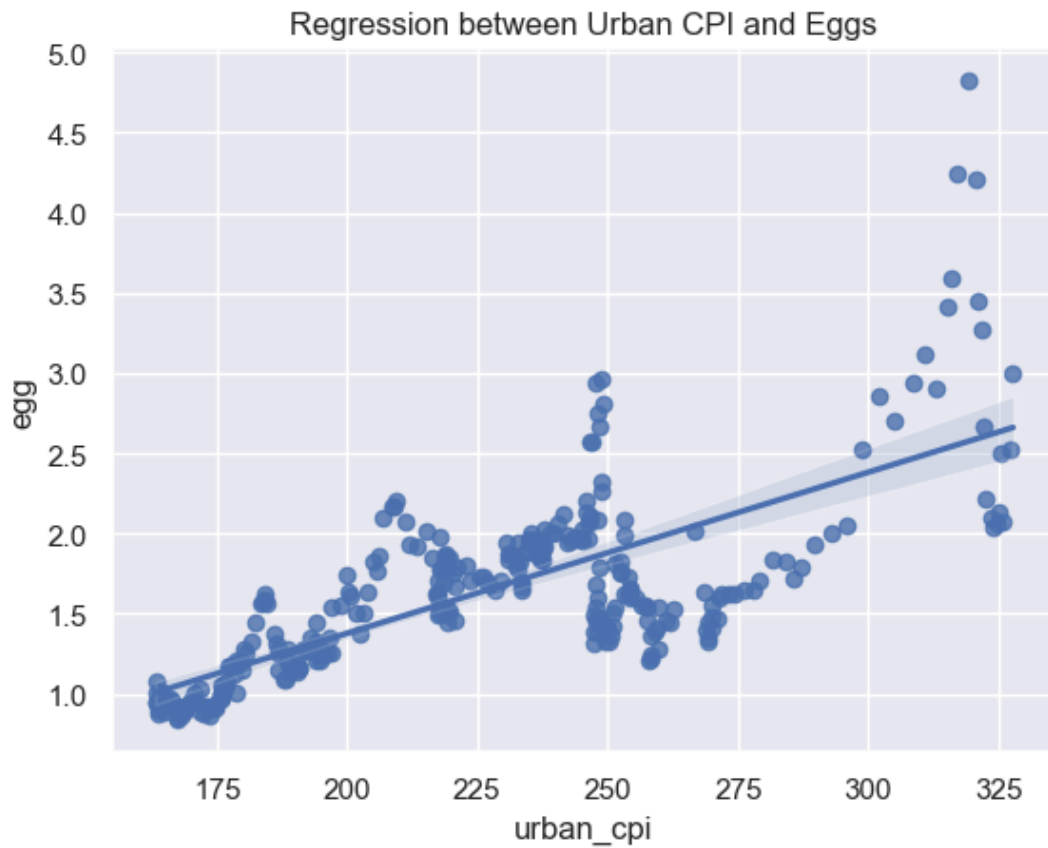
```
In [21]: 1 lr_preds_egg.axes
```

```
Out[21]: [Index([223, 150, 226, 296, 52, 261, 246, 166, 221, 215, 225, 133, 214, 12,
           15, 66, 285, 89, 176, 101, 22, 205, 224, 212, 190, 129, 175, 136,
           152, 55, 201, 253, 144, 8, 271, 111, 263, 81, 184, 282, 5, 59,
           234, 241, 34, 250, 182, 122, 254, 7, 45, 216, 238, 63, 108, 118,
           137, 227, 26, 106, 92],
          dtype='int64')]
```

```
In [22]: 1 plt.scatter(lr_preds_egg.axes, lr_preds_egg)
          2 # plt.title("Regression between Urban CPI and Eggs")
          3 plt.show()
```



```
In [23]: 1 sns.regplot(x='urban_cpi', y='egg', data=groceries_df)
2         plt.title("Regression between Urban CPI and Eggs")
3         plt.show()
```



```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [24]: 1 lr_model_milk = smf.ols(formula="urban_cpi ~ whole_milk", data=X_train).fit()
        2 lr_model_milk.summary()
```

Out[24]: OLS Regression Results

Dep. Variable:	urban_cpi	R-squared:	0.524
Model:	OLS	Adj. R-squared:	0.522
Method:	Least Squares	F-statistic:	262.1
Date:	Tue, 30 Apr 2024	Prob (F-statistic):	3.00e-40
Time:	10:34:50	Log-Likelihood:	-1154.1
No. Observations:	240	AIC:	2312.
Df Residuals:	238	BIC:	2319.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-30.8142	16.054	-1.919	0.056	-62.441	0.813
whole_milk	78.3447	4.839	16.189	0.000	68.811	87.878

Omnibus:	4.496	Durbin-Watson:	2.101
Prob(Omnibus):	0.106	Jarque-Bera (JB):	3.461
Skew:	0.172	Prob(JB):	0.177
Kurtosis:	2.523	Cond. No.	30.2

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

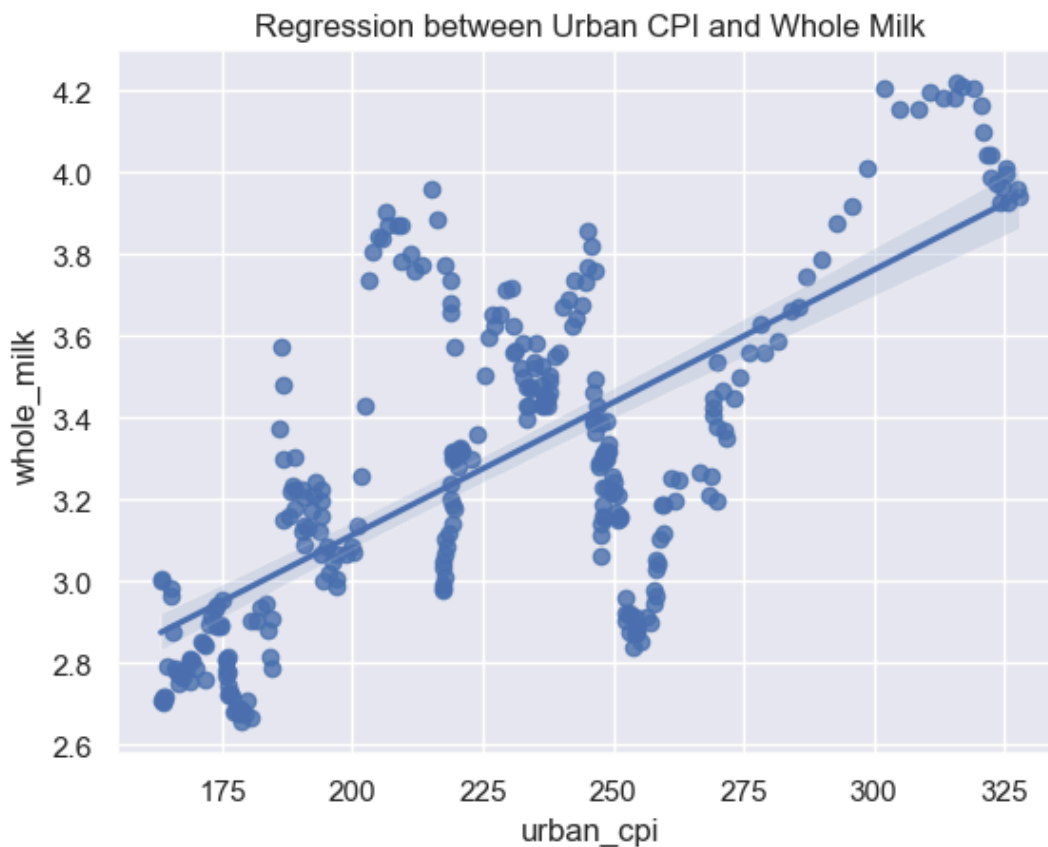
```
In [25]: 1 lr_preds_milk = lr_model_milk.predict(X_test)
        2 lr_preds_milk
```

Out[25]: 223 220.593818
150 260.079533
226 216.363206
296 276.845293
52 178.836108
...
137 228.741664
227 201.164339
26 192.076357
106 272.379647
92 209.233841
Length: 61, dtype: float64

```
In [26]: 1 lr_model_milk.fittedvalues
```

```
Out[26]: 74      221.847333
239      197.403795
64       249.189624
294      276.845293
286      299.095180
...
251      224.040984
192      243.078739
117      261.724771
47       179.619554
172      240.101642
Length: 240, dtype: float64
```

```
In [27]: 1 sns.regplot(x='urban_cpi', y='whole_milk', data=groceries_df)
2 plt.title("Regression between Urban CPI and Whole Milk")
3 plt.show()
```



```
In [28]: 1 # Single LR model between Urban CPI and Banana price
2 lr_model_banana = smf.ols(formula="urban_cpi ~ banana", data=X_train).fit()
3 lr_model_banana.summary()
```

Out[28]: OLS Regression Results

Dep. Variable:	urban_cpi	R-squared:	0.648
Model:	OLS	Adj. R-squared:	0.646
Method:	Least Squares	F-statistic:	437.9
Date:	Tue, 30 Apr 2024	Prob (F-statistic):	7.36e-56
Time:	10:34:51	Log-Likelihood:	-1118.0
No. Observations:	240	AIC:	2240.
Df Residuals:	238	BIC:	2247.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-158.8139	18.523	-8.574	0.000	-195.303	-122.325
banana	686.1932	32.793	20.925	0.000	621.591	750.795

Omnibus:	1.096	Durbin-Watson:	1.960
Prob(Omnibus):	0.578	Jarque-Bera (JB):	1.202
Skew:	-0.130	Prob(JB):	0.548
Kurtosis:	2.771	Cond. No.	26.1

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

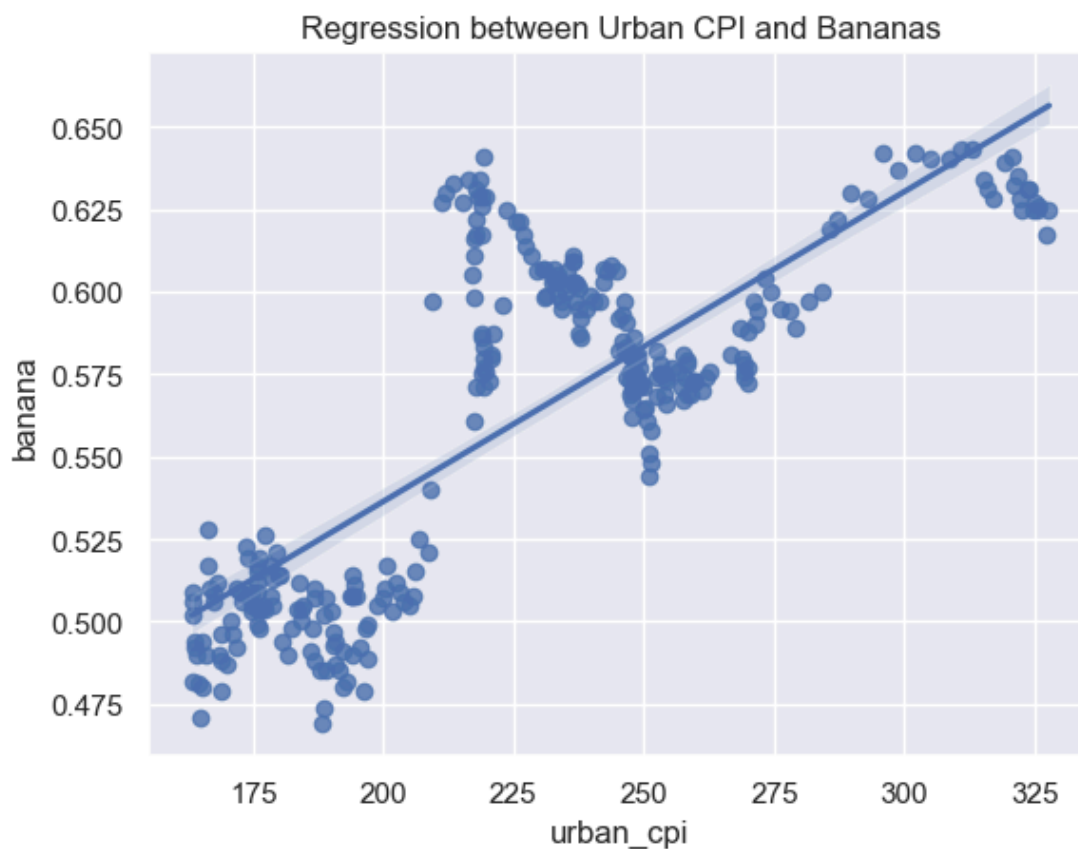
```
In [29]: 1 lr_preds_banana = lr_model_banana.predict(X_test)
2 lr_preds_banana
```

Out[29]: 223 214.475165
150 257.019143
226 224.081870
296 270.743006
52 198.692722
...
137 241.236699
227 230.943801
26 178.793119
106 201.437494
92 176.734539
Length: 61, dtype: float64

```
In [30]: 1 lr_model_banana.fittedvalues
```

```
Out[30]: 74      186.341244
        239      236.433347
        64      182.910278
        294      274.173972
        286      272.115393
        ...
        251      232.316188
        192      246.726245
        117      270.743006
        47       202.123687
        172      254.960563
Length: 240, dtype: float64
```

```
In [31]: 1 sns.regplot(x='urban_cpi', y='banana', data=groceries_df)
        2 plt.title("Regression between Urban CPI and Bananas")
        3 plt.show()
```



```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

```
In [ ]: 1
```

Multiple Linear Regression

To do multiple linear regression on more than one attribute of `groceries_df`, I will pick out specific columns of what is being listed as the independent variable and which columns are considered dependent variables.

I will continue to use `smf.ols` to fit the data and predict.

The columns that are to be used under X: `month_number`, `white_bread`, `whole_milk`, `egg` The column(s) under y: `urban_cpi`

Implement a multilinear regression model on four of the five items and fit the data to the `X_train` dataset.


```
In [32]: 1 # Four grocery items compared: white bread, whole milk, egg, ground beef to be fit
2 lr_model_multi_four = smf.ols('urban_cpi ~ white_bread \
3                               + whole_milk + egg + \
4                               ground_beef', data=X_train).fit()
5 lr_model_multi_four.summary()
```

Out[32]: OLS Regression Results

Dep. Variable:	urban_cpi	R-squared:	0.966
Model:	OLS	Adj. R-squared:	0.966
Method:	Least Squares	F-statistic:	1693.
Date:	Tue, 30 Apr 2024	Prob (F-statistic):	6.51e-172
Time:	10:34:52	Log-Likelihood:	-835.81
No. Observations:	240	AIC:	1682.
Df Residuals:	235	BIC:	1699.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	69.8511	5.213	13.399	0.000	59.580	80.122
white_bread	87.6140	4.814	18.199	0.000	78.130	97.098
whole_milk	-5.4505	2.341	-2.328	0.021	-10.063	-0.838
egg	-1.1216	1.517	-0.739	0.460	-4.111	1.867
ground_beef	21.1033	0.861	24.515	0.000	19.407	22.799

Omnibus:	137.270	Durbin-Watson:	2.259
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2311.190
Skew:	1.841	Prob(JB):	0.00
Kurtosis:	17.750	Cond. No.	56.2

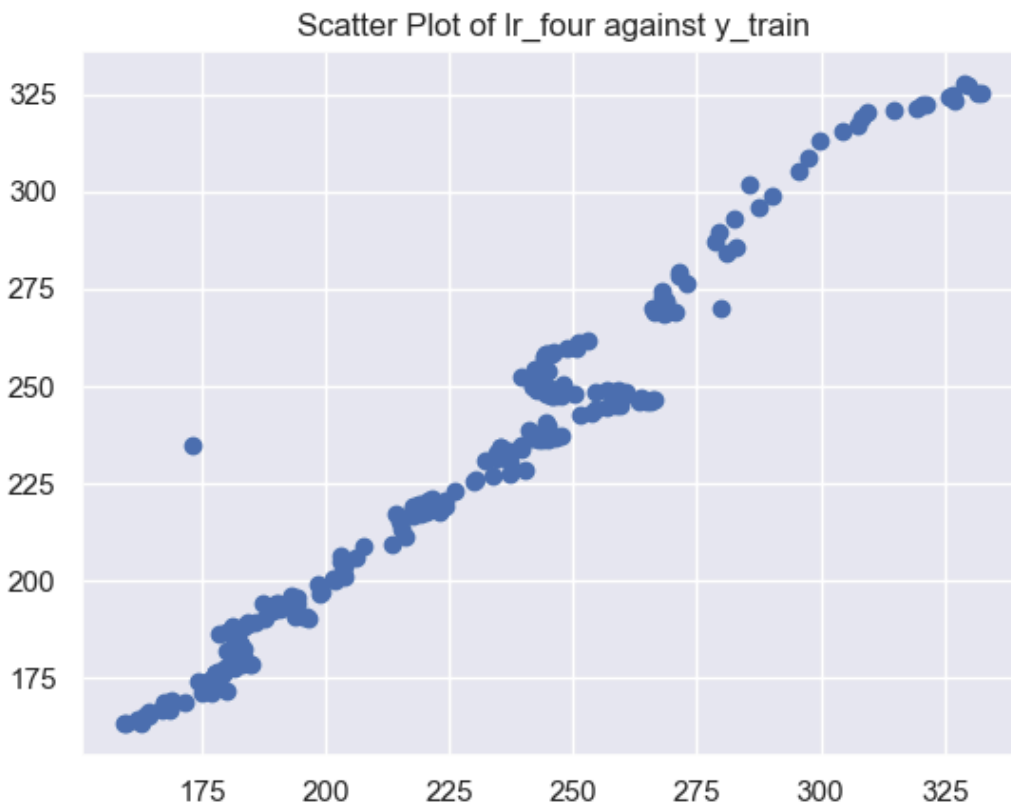
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

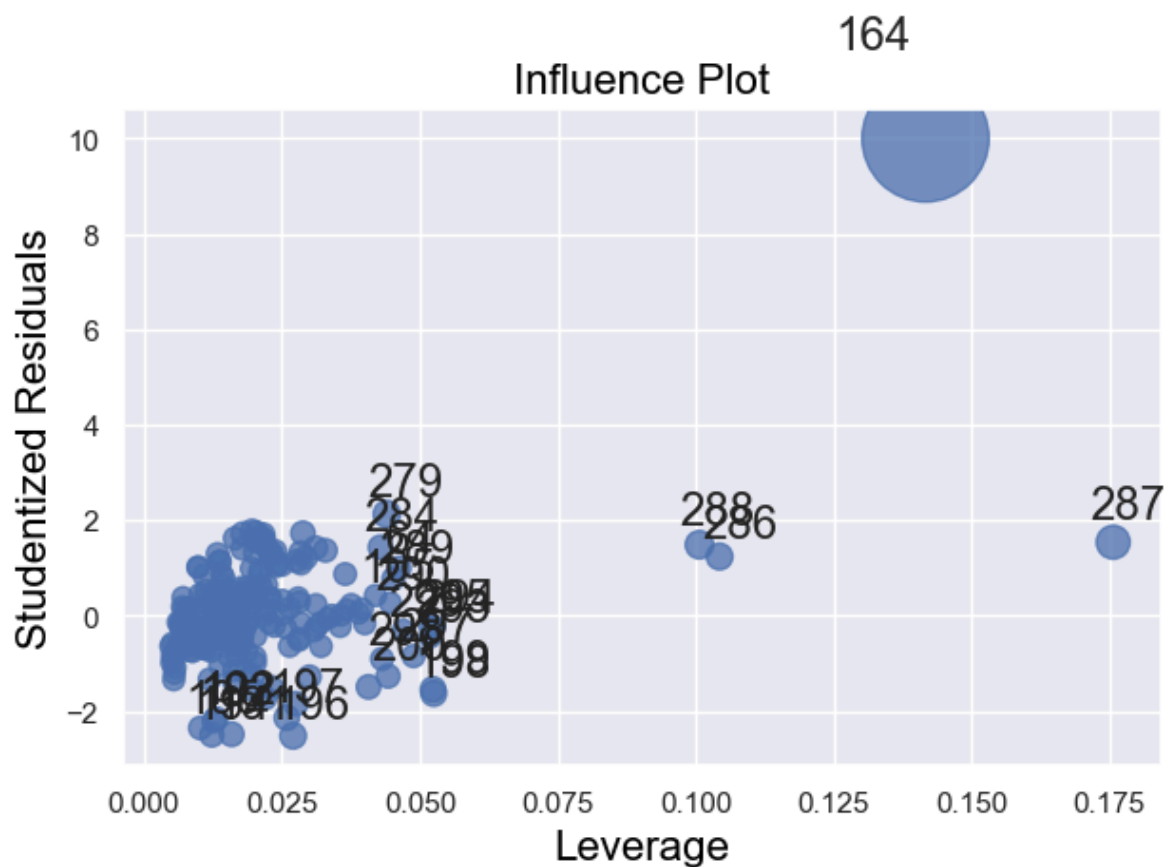
```
In [33]: 1 # Predict the CPI values (four items)
        2 lr_four = lr_model_multi_four.predict(X_train)
        3 lr_four
```

```
Out[33]: 74      187.659707
        239      244.064792
        64      178.527382
        294      325.875595
        286      307.529190
        ...
        251      250.855988
        192      263.616216
        117      217.822411
        47       181.526868
        172      246.366184
        Length: 240, dtype: float64
```

```
In [62]: 1 # Scatter plot of predicted results against actual CPI values
        2
        3 plt.scatter(lr_four, y_train)
        4 plt.title("Scatter Plot of lr_four against y_train")
        5 plt.show()
```



```
In [36]: 1 # Influence plot on the four items
2
3 fig = sm.graphics.influence_plot(lr_model_multi_four)
4 fig.tight_layout(pad=1.0)
5 plt.show()
```

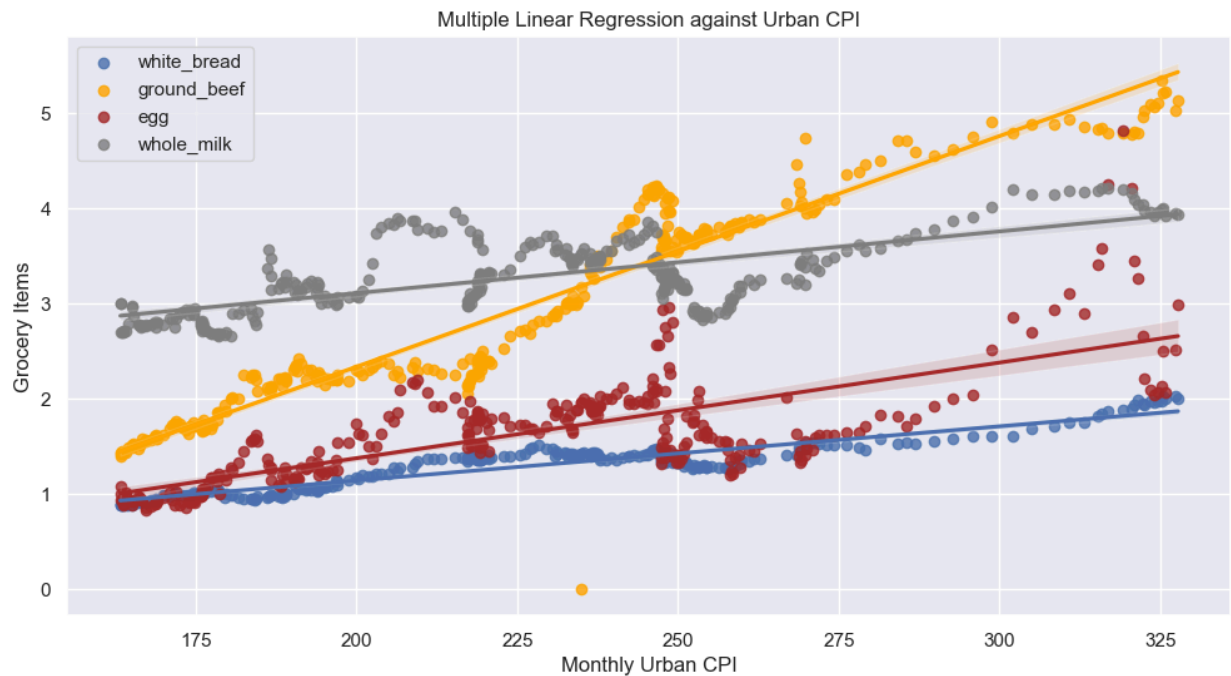


In [63]:

```

1 # Four Grocery Items' Regression Plots
2
3 plt.figure(figsize=(12,6))
4 sns.regplot(x='urban_cpi', y='white_bread', data=groceries_df, label='white_bread')
5 sns.regplot(x='urban_cpi', y='ground_beef', data=groceries_df, label='ground_beef')
6 sns.regplot(x='urban_cpi', y='egg', data=groceries_df, label='egg', color='brown')
7 sns.regplot(x='urban_cpi', y='whole_milk', data=groceries_df, label='whole_milk',
8
9 plt.title("Multiple Linear Regression against Urban CPI")
10 plt.xlabel("Monthly Urban CPI")
11 plt.ylabel("Grocery Items")
12 plt.legend()
13 plt.show()

```



```
In [64]: 1 # All five grocery items compared: white bread, whole milk, egg, ground beef, banana
2
3 lr_model_multi_five = smf.ols('urban_cpi ~ white_bread \
4                               + whole_milk + egg + \
5                               ground_beef', data=X_train).fit()
6 lr_model_multi_five.summary()
```

Out[64]: OLS Regression Results

Dep. Variable:	urban_cpi	R-squared:	0.966
Model:	OLS	Adj. R-squared:	0.966
Method:	Least Squares	F-statistic:	1693.
Date:	Tue, 30 Apr 2024	Prob (F-statistic):	6.51e-172
Time:	10:52:05	Log-Likelihood:	-835.81
No. Observations:	240	AIC:	1682.
Df Residuals:	235	BIC:	1699.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	69.8511	5.213	13.399	0.000	59.580	80.122
white_bread	87.6140	4.814	18.199	0.000	78.130	97.098
whole_milk	-5.4505	2.341	-2.328	0.021	-10.063	-0.838
egg	-1.1216	1.517	-0.739	0.460	-4.111	1.867
ground_beef	21.1033	0.861	24.515	0.000	19.407	22.799

Omnibus:	137.270	Durbin-Watson:	2.259
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2311.190
Skew:	1.841	Prob(JB):	0.00
Kurtosis:	17.750	Cond. No.	56.2

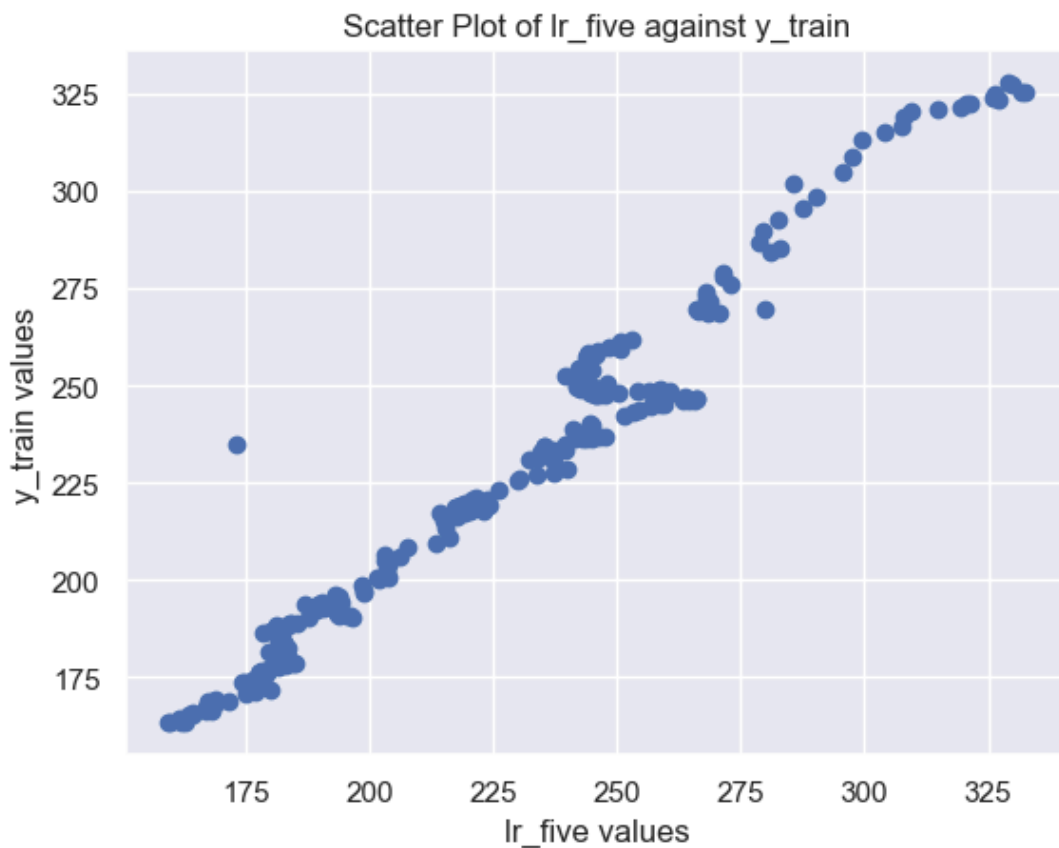
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [61]: 1 # Predict the CPI values (four items)
2
3 lr_five = lr_model_multi_five.predict(X_train)
4 lr_five
```

```
Out[61]: 74      187.659707
239      244.064792
64       178.527382
294      325.875595
286      307.529190
...
251      250.855988
192      263.616216
117      217.822411
47       181.526868
172      246.366184
Length: 240, dtype: float64
```

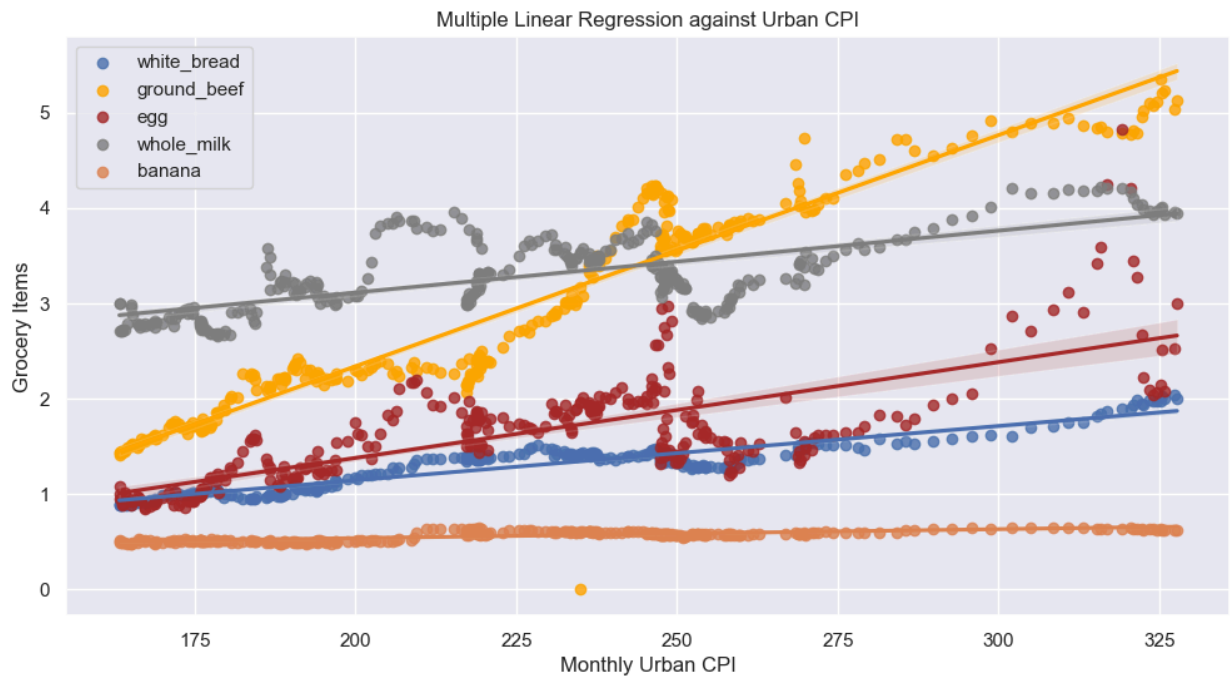
```
In [65]: 1 # Scatter plot of predicted results on the five items against actual CPI values
2
3 plt.scatter(lr_five, y_train)
4 plt.title("Scatter Plot of lr_five against y_train")
5 plt.xlabel("lr_five values")
6 plt.ylabel("y_train values")
7 plt.show()
```



```

In [59]: 1 # Five Grocery Items' Regression Plots
2
3 plt.figure(figsize=(12,6))
4 sns.regplot(x='urban_cpi', y='white_bread', data=groceries_df, label='white_bread')
5 sns.regplot(x='urban_cpi', y='ground_beef', data=groceries_df, label='ground_beef')
6 sns.regplot(x='urban_cpi', y='egg', data=groceries_df, label='egg', color='brown')
7 sns.regplot(x='urban_cpi', y='whole_milk', data=groceries_df, label='whole_milk',
8             sns.regplot(x='urban_cpi', y='banana', data=groceries_df, label='banana')
9
10 plt.title("Multiple Linear Regression against Urban CPI")
11 plt.xlabel("Monthly Urban CPI")
12 plt.ylabel("Grocery Items")
13 plt.legend()
14 plt.show()

```



```

In [ ]: 1

```

We can see that based on the R-squared value that there is a positive relationship between the Urban CPI and any combination of four of the five grocery items. In this case, I compared the Urban CPI to the sum of the coefficients of white bread, whole milk, eggs, and ground beef.

```

In [ ]: 1

```

```
In [37]: 1 # Just using two grocery items for comparison: whole milk and egg
          2
          3 lr_model_multi_two = smf.ols('urban_cpi ~ whole_milk + egg', data=X_train).fit()
          4 lr_model_multi_two.summary()
```

Out[37]: OLS Regression Results

Dep. Variable:	urban_cpi	R-squared:	0.610
Model:	OLS	Adj. R-squared:	0.607
Method:	Least Squares	F-statistic:	185.2
Date:	Tue, 30 Apr 2024	Prob (F-statistic):	3.70e-49
Time:	10:34:55	Log-Likelihood:	-1130.3
No. Observations:	240	AIC:	2267.
Df Residuals:	237	BIC:	2277.
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	38.5488	17.454	2.209	0.028	4.164	72.933
whole_milk	40.7878	6.810	5.989	0.000	27.372	54.204
egg	32.7281	4.536	7.215	0.000	23.792	41.664

Omnibus:	15.872	Durbin-Watson:	2.124
Prob(Omnibus):	0.000	Jarque-Bera (JB):	15.616
Skew:	0.574	Prob(JB):	0.000406
Kurtosis:	2.505	Cond. No.	41.7

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

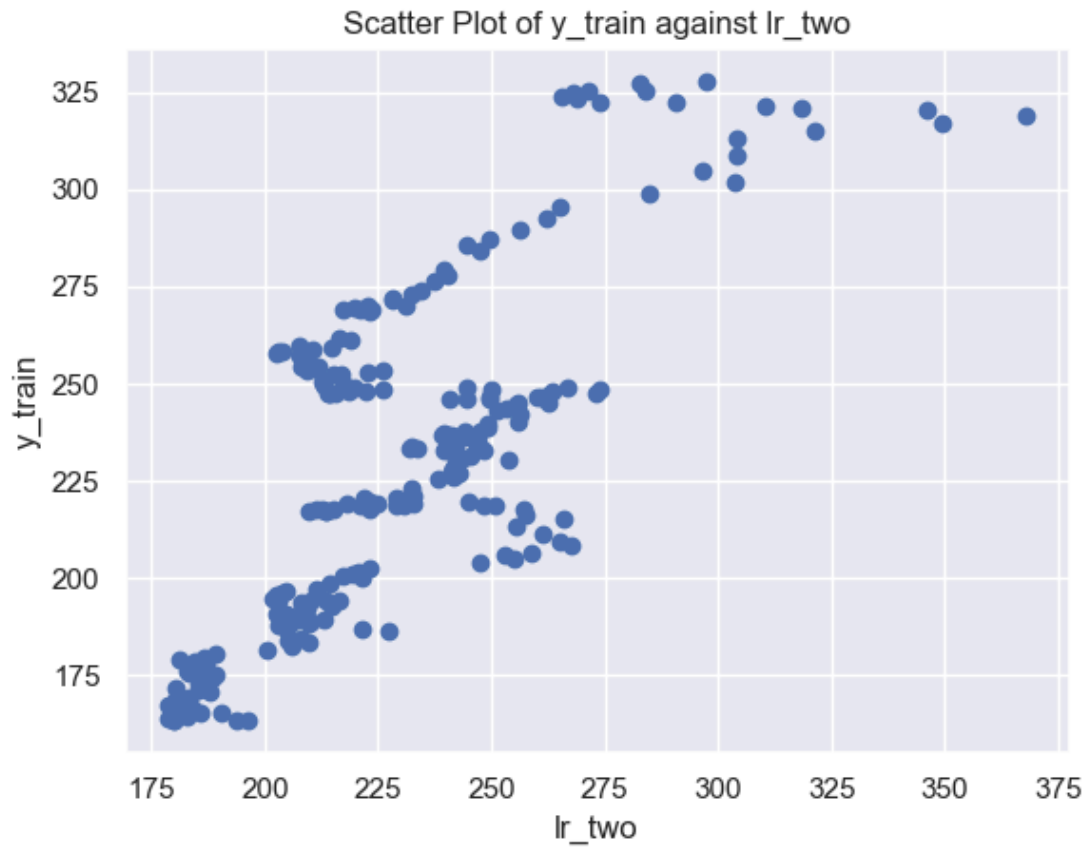
```
In [38]: 1 lr_two = lr_model_multi_two.predict(X_train)
          2 lr_two
```

Out[38]:

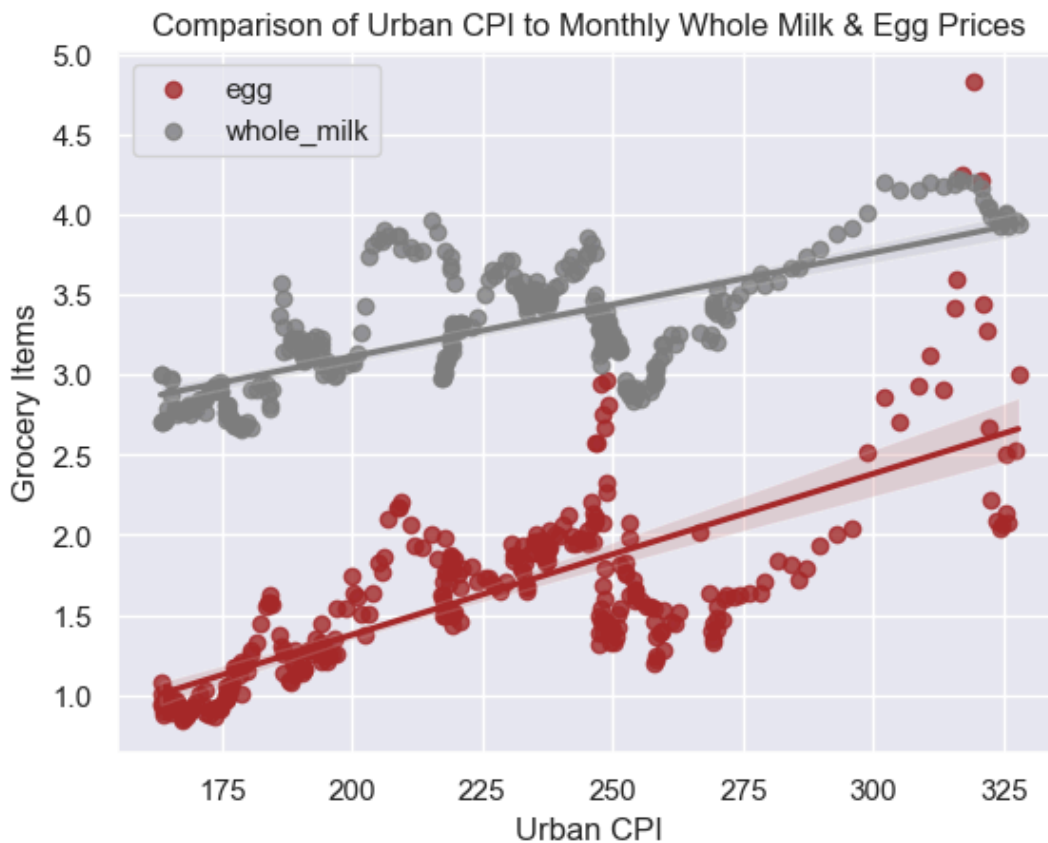
```
74      208.185030
239      208.223163
64       227.231012
294      265.586057
286      349.400619
...
251      219.047322
192      249.479258
117      251.004752
47       186.560390
172      240.434595
Length: 240, dtype: float64
```



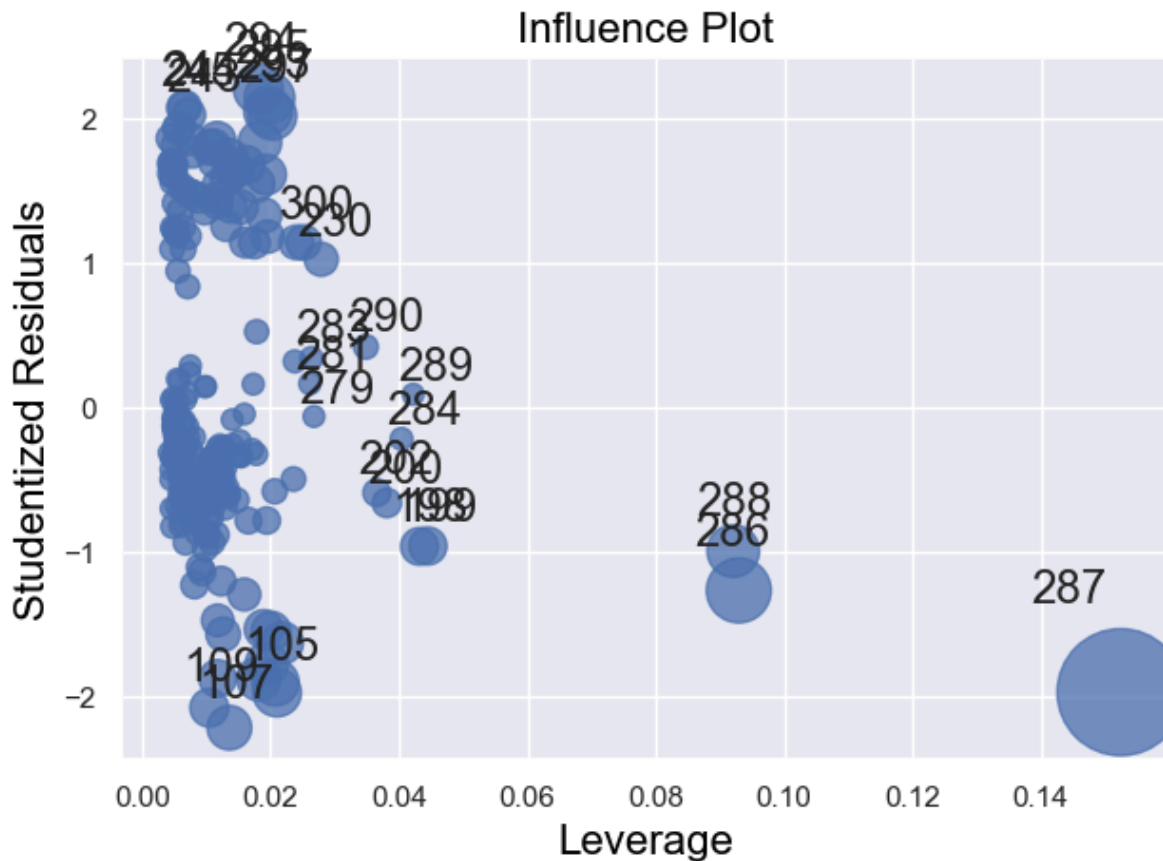
```
In [67]: 1 # Note how the scatter plot for two items is different from those of four and five
2 # The regression line would be off
3
4 plt.scatter(lr_two, y_train)
5 plt.title("Scatter Plot of y_train against lr_two")
6 plt.xlabel("lr_two")
7 plt.ylabel("y_train")
8 plt.show()
```



```
In [68]: 1 # Regression plots of egg and whole milk
2 sns.regplot(x='urban_cpi', y='egg', data=groceries_df, label='egg', color='brown')
3 sns.regplot(x='urban_cpi', y='whole_milk', data=groceries_df, label='whole_milk',
4 plt.ylabel("Grocery Items")
5 plt.xlabel("Urban CPI")
6 plt.title("Comparison of Urban CPI to Monthly Whole Milk & Egg Prices")
7 plt.legend()
8 plt.show()
```



```
In [41]: 1 fig = sm.graphics.influence_plot(lr_model_multi_two)
2 fig.tight_layout(pad=1.0)
3 plt.show()
```



```
In [ ]: 1
```

```
In [ ]: 1
```

Conclusion

We can see that all of the five grocery items picked for this evaluation have had their average prices increased over the 25-year period. Each of the five individual grocery items, when compared to the monthly urban CPI, had some form of correlation to the urban CPI. Of the five, however, the item which had data points that were closest to the regression coefficients was ground beef, where the monthly average price by the pound was nearly steady with the rise in the CPI.

The comparison of two different kinds of linear regression models from two different Python libraries, with two different methods of conducting this form of supervised learning. The statsmodels version provided immediate fitting and prediction value sets that could easily be presented visually, and provides entire arrays worth of predicted values with respect to the urban CPI.

It can be seen that the fewer items are involved in prediction models, the less likely that one item will have a strong connection with the CPI. If all five items are included, the connection is higher. Both single and multiple linear regression show that this will be the case with the prices of items gradually rising despite periods of faster price hikes.

These five grocery items alone only portray just a small picture of how it can be affected by the overall urban CPI, and the CPI itself. The overall cost of groceries have gone up over time, independent of other external factors (public health crises, climate change, rapid economic change, etc.). Even looking at just the average monthly prices at face value, without the analysis, prices have been and will continue to go up. Supply and demand might be one of the key causes of the increase in costs and spending, but how tightly related produce, meats, or grains, among other agricultural fields, are to CPI here in the US and around the world, will depend on other factors in how consumers purchase what food they need.

The price of items as well as the CPI in the US will rise gradually and not accelerate as initially thought. Despite challenges with being able to explicitly outline the constant rise, there will be no faster growth. When prices of grocery items drop or remain the same on a monthly basis, the mean price growth will slow down and the slope be closer to zero. Barring external factors and the types of items, growth of the CPI will likely continue to increase while the rate of the prices of items purchased will increase at slower,

In []:

1

References

FRED, Federal Reserve Bank of St. Louis. (n.d.). *Average Price: Bananas (Cost per Pound/453.6Grams) in U.S. City Average [APU0000711211]*. [Unpublished raw data]. Retrieved April 3, 2024, from <https://fred.stlouisfed.org/series/APU0000711211>.

FRED, Federal Reserve Bank of St. Louis. (n.d.). *Average Price: Bread, White, Pan (Cost perPound/453.6 Grams) in U.S. City Average [APU0000702111]*. [Unpublished raw data]. Retrieved April 2, 2024, from <https://fred.stlouisfed.org/series/APU0000702111> (<https://fred.stlouisfed.org/series/APU0000702111>).

FRED, Federal Reserve Bank of St. Louis. (n.d.). *Average Price: Eggs, Grade A, Large (Cost per Dozen)in U.S. City Average [APU0000708111]*. [Unpublished raw data]. Retrieved March 27, 2024, from <https://fred.stlouisfed.org/series/APU0000708111> (<https://fred.stlouisfed.org/series/APU0000708111>).

FRED, Federal Reserve Bank of St. Louis. (n.d.). *Average Price: Ground Beef, 100% Beef (Cost perPound/453.6 Grams) in U.S. City Average [APU0000703112]*. [Unpublished raw data]. Retrieved April 8, 2024, from <https://fred.stlouisfed.org/series/APU0000703112> (<https://fred.stlouisfed.org/series/APU0000703112>).

FRED, Federal Reserve Bank of St. Louis. (n.d.). *Consumer Price Index for All Urban Consumers: Foodin U.S. City Average [CPIUFDNS]*. [Unpublished raw data]. Retrieved March 31, 2024, from <https://fred.stlouisfed.org/series/CPIUFDNS> (<https://fred.stlouisfed.org/series/CPIUFDNS>).

FRED, Federal Reserve Bank of St. Louis. (n.d.). *Producer Price Index by Commodity: Processed Foods and Feeds: Fluid Whole Milk [WPU02310301]*. [Unpublished raw data]. Retrieved April 28, 2024, from <https://fred.stlouisfed.org/series/WPU02310301> (<https://fred.stlouisfed.org/series/WPU02310301>).

In []:

1

In []:

1

